

IP/NDN Gateway and Bridge to Support Heterogeneous Network Interoperability

Christopher A. Wood^{*}
University of California Irvine
Irvine, CA, USA
woodc1@uci.edu

ABSTRACT

TODO

Categories and Subject Descriptors

H.4 [TODO]: TODO; D.2.8 [TODO]: TODO—*TODO*

General Terms

TODO

Keywords

TODO

1. INTRODUCTION

At its core, the design and architecture of today's Internet is communication-based packet switching network. Since its inception, it has been retrofitted with a variety of transport and application layer protocols and middleware to support a growing set of consumer applications, such as the Web, email, and perhaps most importantly in recent times, media streaming services. The latter type of applications are bandwidth-intensive content distribution applications which leverage the underlying communication-based network as a distribution network, leading to a massive consumption of vital networking resources.

Information-centric networks (ICNs) are a new class of network architecture designs that aim to address this increasingly popular type of network traffic by decoupling data from its source and shifting the emphasis of addressable content from hosts and interfaces to content [1]. By directly addressing content instead of hosts, content dissemination and security can be “distributed” throughout the network in the sense that consumers requests for content are satisfied by *any* resource in the network (i.e., not necessarily the original producer). For example, routers close to consumers may cache content with a particular name and then satisfy all content requests that match the content's name. In-network caching and data-centric security measures are two of

the defining characteristics of these new network designs.

As of today there are several information-centric networking proposals being explored as alternative designs to today's Internet; Named Data Networking (NDN) [2] is one of the more promising designs that is still an active area of active research (see www.named-data.net for more information). As a replacement for IP-based networks, the complete adoption of NDN, or any one of these designs, will realistically need to be done by slow and continual integration and replacement of IP-based networking resources with NDN-based resources. Currently, however, there is no engineering plan to support the IP-to-ICN integration without significant software modification and application, transport, or network layer source code modifications.

Consequently, the primary objective of this work is to aid the integration of future content-centric networking resources into the existing IP-centric Internet by providing an application and transport layer gateway between IP and NDN resources. We assume that NDN *will not* be deployed over the IP network, as such a deployment scheme would nullify the need for a gateway between the two networks. Application-layer traffic corresponding to protocols such as HTTP, FTP, SMTP, IMAP, etc. will be translated by middleware running in such gateways to correctly interface with the NDN resources, thereby serving as a semantic bridge between these two fundamentally different networking architectures. Similarly, using a custom NDN-to-IP protocol, NDN *interests* for content will be translated to the messages adhering to the corresponding application-layer protocol to traverse the IP network. When serving as a bridge between islands of NDN resources, two gateways will leverage the features of the IP network to move messages among two separate NDN beds. Specifically, NDN interests may be encapsulated in UDP datagrams composed of IP packets traversing an IP network between two NDN islands. The primary benefit of this semi-transparent gateway is that existing IP-centric applications need not be modified at any layer in the network stack to interoperate with NDN resources.

^{*}NSF GRFP blurb.

2. HETEROGENEOUS NETWORKS AND THE NEED FOR A GATEWAY AND BRIDGE

Consider the typical hourglass network stack in IP-based networks as shown in left-hand image of Figure 2. This layered design with a thin-waist infrastructure (IP packets for traffic flow) is what enabled the Internet to grow and expand at such a rapid rate; higher layers in the protocol stack extend this communication medium with support for a variety of applications and networking features (e.g., reliable message traversal via TCP). While the NDN architecture introduces a fundamental paradigm shift in the way information is published and retrieved on a network, its design, shown at a high level in the right-hand image of Figure 2, borrows the same hourglass design as IP networks. Observe that upper layers of the network stack still promote the development of robust applications based on the underlying communication layers. The difference, however, is that network traffic flow management (i.e., to enable reliable and stable communication) and security are *built into* the network stack. These architectural differences mean that application, transport, and network layer protocol semantics in IP-based networks are distinct from protocol semantics in NDN networks. The NDN gateway is intended to bridge between IP and NDN networks by performing this semantic translation between protocols.




Figure 1: A visual comparison of the network stacks for the IP and NDN network architectures.

Assuming that NDN is not to be deployed over IP, but instead as a separate network stack entirely, the need for such a gateway cannot be understated. Consider two instances of an application, A and B , that wish to send data back and forth to each other. Application A is running on a host with only an IP interface, and application B is running on a host with only an NDN interface. What does it mean for application A to establish a TCP connection stream with application B and what does it mean for application B to issue an interest to application A when neither application speaks the network language of the other. In order for these

two applications to communicate, the semantics of a TCP stream-oriented connection must be translated to a stream of contiguous interests, and vice versa.

Now consider an alternative scenario in which two “islands” of NDN hosts exist in isolation, both of which running the same instance of an application. Let host H_1 and H_2 be a host in the first and second island, respectively, that wish to retrieve data from one another using interests. Without a network route between the two islands, there would be no way for H_1 and H_2 to interact with one another. However, if there existed two gateways at the borders of each island, each of which connected to the same IP network, then interests from H_1 (H_2) could be sent to H_2 (H_1) as follows:

1. An interest from H_1 is intercepted a gateway G_1 .
2. G_1 encapsulates the interest in an IP packet sent to gateway G_2 adhering to IPsec.
3. G_2 unwraps and re-issues the interest and waits for the content to be retrieved from H_2 . Upon reception, the content’s signature is verified and the content is encapsulated in a signed IPsec packet and sent to G_1 .
4. G_1 verifies the signature of the packet, creates and signs a new content object, and sends the content object back downstream to H_1 .

3. NETWORK SEMANTIC TRANSLATIONS AT THE GATEWAY

Traditional IP-based applications and upcoming NDN-based applications treat both the network and content in significantly different ways. In IP-based settings, application and transport layer mechanisms and protocols leverage the underlying IP network layer to send packets to specific hosts. In contrast, in NDN-based settings, there are no straightforward application or transport layer analogs; the network layer, responsible for the issuance of interests and retrieval of content, is abstracted to authenticated (i.e., digitally signed) content objects or streams of data that are consumed by applications. From this perspective, there is no clear bijection between application and transport layer communication in IP-based settings and content and stream-centric data retrieval in NDN-settings. Therefore, to support the interoperability of these two networking paradigms, we need a mechanism to translate the semantics of IP-based communication to and from NDN-based content retrieval.

The direction of traffic across the gateway has a strong influence on how this semantic translation is done: IP-to-NDN application and transport layer protocols will be mapped to corresponding NDN-interests, and NDN-to-IP traffic (in the form of interests) will be encoded so as to map to the appropriate IP-based application

or transport layer protocol. Stateless protocols such as HTTP greatly simplify the job of the gateway because it need not maintain any state to support communication across both networks. However, stateful protocols, such as TCP, naturally require the gateway to maintain state so as to emulate the behavior of an endpoint or host that implements such protocols. For instance, if an IP-based application wishes to establish a TCP connection with an application running on an NDN network to retrieve data, then the gateway must maintain stateful information needed to transform streams of data retrieved over a TCP socket to (packed) discrete and contiguous interests in the NDN network. In what follows we describe the semantic translation details for application and transport layer stateful and stateless protocols in both directions across the gateway.

3.1 NDN-to-IP Application Layer Semantic Translation

There does not exist a native correspondence between NDN interests and IP-based application layer protocol messages. For example, there is no standard way for a client to represent an HTTP GET request in the format of an NDN interest. To make this type of semantic translation possible, the NDN-to-IP application layer bridge leverages the human-readable names of content to encode IP-based application layer protocol specifics. The encoding for HTTP and FTP semantic translations is specified in EBNF form below; other application-layer protocols can easily be supported by extending this grammar in the natural way.

Interests encoded using this grammar are intercepted in the `NDNInputStage` of the gateway (see Figure 3). Upon reception, the gateway parses the message, stores a new entry in the NDN pending message table, and forwards the decoded message contents to the appropriate `IPOutputStage` pipeline stage. Upon reception, the IP response is retrieved in the `IPInputStage` and forwarded inwards to the `NDNOutputStage`, where the corresponding entry in the NDN pending message table is indexed using the contents of the arriving message to retrieve the original incoming interest name. Once fetched, the gateway creates and signs a new content object with the IP network response, and then forwards the content downstream to its intended consumer.

Unlike traditional NDN routing, the NDN pending message table does not collapse interests by default. The reason for this is that application-protocol interests are often issued when *new* state needs to change (i.e., cached responses not generated on demand from the intended IP recipient are not acceptable).

It is important to note that stateful protocols such as FTP must explicitly embed identifying information about the consumer in order to operate correctly. The reason is that one FTP session must be associate with

at most one consumer, and since NDN does not have any notion of host addresses or identifiers, the consumer application must explicitly encode its identity in the interest. Our grammar proposes embedding the consumer public key and a corresponding digital signature of the entire interest for such protocols so that consumers can be explicitly identified and their sessions cannot be hijacked by other consumers (doing so would require compromising a consumer and its private key).

4. BRIDGING ISOLATED NETWORKS VIA MESSAGE ENCAPSULATION AT THE BRIDGE

TODO

5. REFERENCES

- [1] V. Jacobson, D. K. Smetters, J. D. Thornton, M. F. Plass, N. H. Briggs, and R. L. Braynard. Networking named content. In *Proceedings of the 5th International Conference on Emerging Networking Experiments and Technologies*, CoNEXT '09, pages 1–12, New York, NY, USA, 2009. ACM.
- [2] L. Zhang, D. Estrin, J. Burke, V. Jacobson, J. D. Thornton, D. K. Smetters, B. Zhang, G. Tsudik, D. Massey, C. Papadopoulos, et al. Named data networking (ndn) project. *Relatório Técnico NDN-0001*, Xerox Palo Alto Research Center-PARC, 2010.

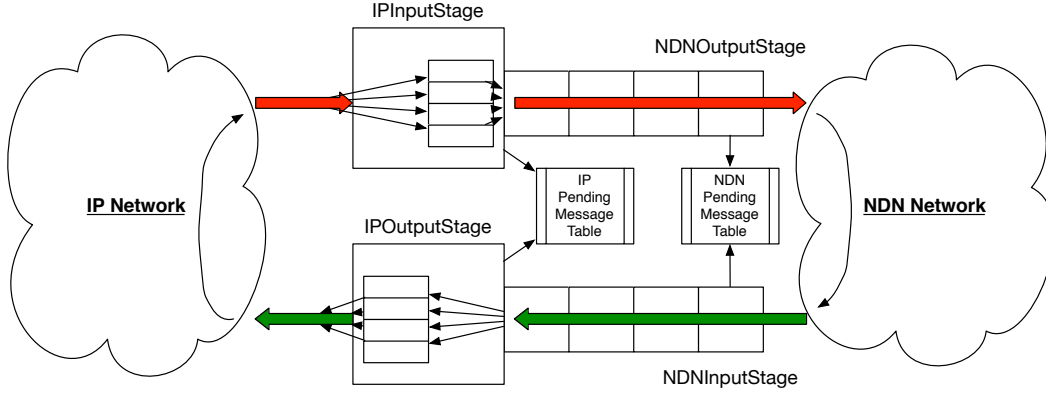


Figure 2: Bidirectional message pipeline for IP-to-NDN and NDN-to-IP message traversal.

```

<ip-interest> → '/.../ip/'<protocol>
<protocol> → 'http/'<http-cmd>['/'<http-path>] | 'ftp/'<ftp-cmd>['/'<ftp-param>] '/'<ftp-ident>
<http-cmd> → 'GET' | 'PUT' | 'POST' | 'DELETE'
<ftp-cmd> → 'ascii' | 'binary' | 'bye' | 'cd' | 'close' | 'delete' | 'get' | 'help' | 'lcd' | 'ls' | 'mkdir' | 'mget' | 'open' |
'put' | 'pwd' | 'quit' | 'rmdir'
<http-path> → <uri> | ip-address'→'port'/'<uri-encoded-string>
<ftp-ident> → <SHA256-hash>'/'<public-key>
<ftp-param> → <uri-encoded-string>
<uri-encoded-string> → TODO

```