

# Multipurpose IP/NDN Gateway and Bridge for Heterogeneous Network Interoperability

Christopher A. Wood  
[www.christopher-wood.com](http://www.christopher-wood.com)  
[woodc1@uci.edu](mailto:woodc1@uci.edu)

May 21, 2014

# Agenda

Overview

Modes of Operation

Gateway Functionality

Bridge Functionality

Internal Design

Experimental Setups

# Today's Internet: Communication Networks as Distribution Networks

The communication-centric design enables point-to-point communication between any two parties:

- ▶ Names and interfaces
- ▶ Supports end-to-end conversations
- ▶ Provides unreliable packet delivery via IP datagrams
- ▶ Compensates for simplicity of IP via complexity of TCP

Important observations:

- ▶ Helped facilitate today's concentric world, *but was never designed for it!*
- ▶ Fundamental communication model: point-to-point conversation between two hosts (IP interfaces)
- ▶ The central abstraction is a host identifier corresponding to an IP address

# Data vs Communication Networks

Distribution/data (DN) and communication (CN) networks differ in several key ways:

	CN	DN
Naming	Endpoints	Content
Memory	Invisible & limited	Explicit (storage = wires)
Security	Communication process/channel	Content

# CCN Overview

Content-centric networking flips around the host-based model of the Internet architecture

- ▶ *Content names*, rather than content locations, become addressable.
- ▶ The network is permitted to store (cache) content that is in high demand
- ▶ End result: less traffic to/from the content's original source, better usage of network resources, less latency, etc etc.

## CCN Overview (continued)

How is data actually retrieved?

- ▶ A consumer  $C$  sends out an *interest* for content they desire.
- ▶ A router  $R_i$  use the information in their forwarding information base (FIB) table and data in cached in their content store (CS) to handle incoming interests:
  1. If content with the same name matches what's stored in the CS, return that content
  2. Else, store the interest in their pending interest table (PIT) (including the downstream router  $R_{i-1}$  or consumer  $C$  that made the request), and forward the request upstream to the next router  $R_{i+1}$  based on their FIB.
  3. FIBs are configured using protocol similar to OSPF
- ▶ Once the interest is satisfied in  $R_i$ , the PIT entry is cleared, the content is cached, and the data is sent downstream to  $C$  or  $R_{i-1}$ .

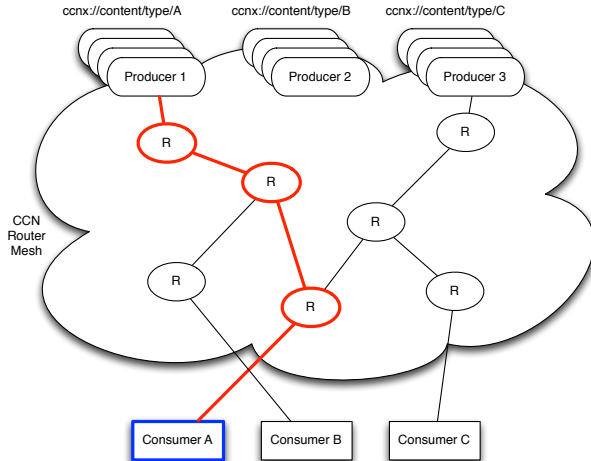
# Interest Format

- ▶ Interests are similar to URLs:

`ccnx://rit/gccis/cs/spr/ramsey_survey`

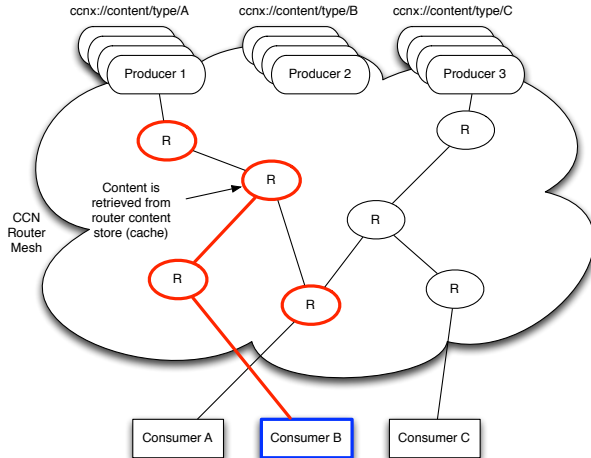
- ▶ The / character is a delimiter that separates name *components*
- ▶ A component can be *anything*, including binary data (e.g. ciphertext)
- ▶ Interests are matched to providers in FIBs using a standard longest-prefix rule (to my knowledge, interests in CSs must match completely)

# CCN in Action - #1

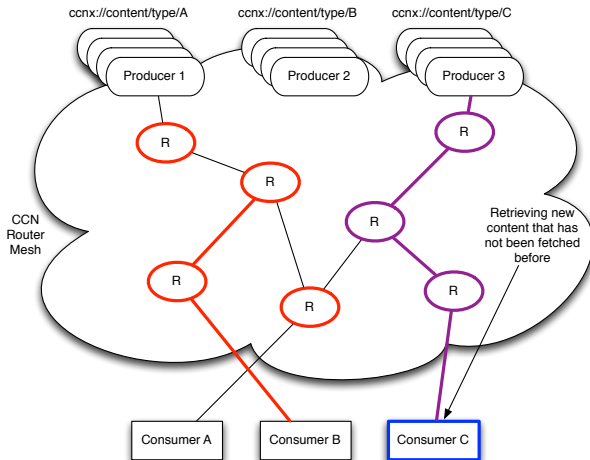




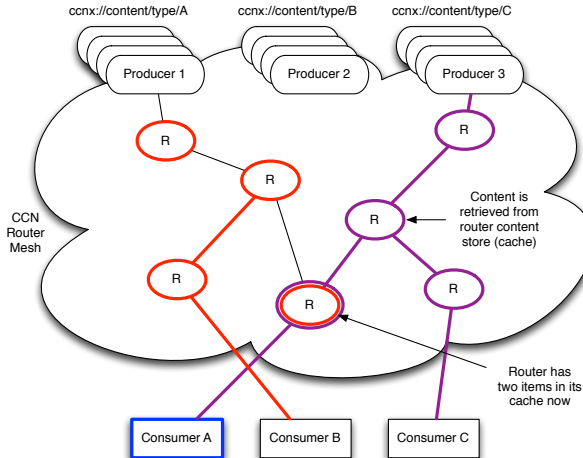
## CCN in Action - #2




## CCN in Action - #3



## CCN in Action - #4



# CCN at a Larger Scale



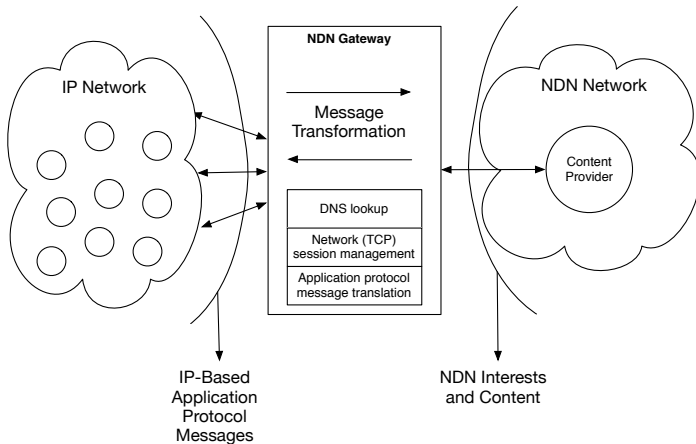
## Motivation for NDN Gateway/Bridge

**Question:** If adopted, how will NDN/CCN be deployed?

1. “Turn off” the Internet, swap in new hardware, and then flip the switch again
  - ▶ Bad idea...
2. Incrementally “roll out” NDN hardware and slowly make it interoperable with existing IP network
  - ▶ How to enable NDN-based applications to communicate with IP-based applications (and vice versa)?
  - ▶ ...and how to do this without re-writing the transport/network layer of IP-based applications to use CCNx (i.e., implement CCN functionality on top of IP)?

**Answer:** Use a NDN-network edge gateways to hide the details of NDN/IP communication mechanics and translate IP messages to compliant NDN interests (and vice versa), and use NDN-network edge bridges to connect isolated NDN “islands”.

# Gateway Semantic Translations



## IP-to-NDN Traffic

- ▶ HTTP GET requests issued to get content with a similar name
  - ▶ e.g., GET X.X.X.X:80/ndn/ccnx/name/of/content
  - ▶ The request path is mapped to the outgoing interest name
- ▶ TCP connections established to stream data to NDN producers
  - ▶ Socket connection between IP-based client and gateway established, NDN producer name first sent, and then all remaining data is streamed
  - ▶ The gateway partitions data from the socket and packs it into an interest for the desired NDN producer


## NDN-to-IP Traffic

- ▶ Interests are encoded according to a special grammar to enable the gateway to parse interests and issue them using the appropriate IP-based protocol

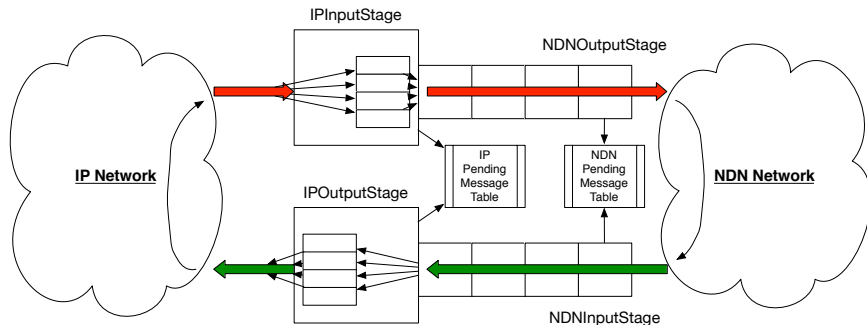
```
<ip-interest>: '/.../ip/'<protocol>.  
<protocol>: 'http/'<http-cmd>['/'<http-path>] | 'tcp/'<tcp-ident>/'<uri-  
encoded-string>.  
<http-cmd>: 'GET' | 'PUT' | 'POST' | 'DELETE'.  
<http-path>: <uri> | <ip-address>[port]['/'<uri-encoded-string>]  
<tcp-ident>: <SHA256-hash>/'/'<nonce>/'/'<public-key>.  
<tcp-param>: <uri-encoded-string>.
```



## Bridging NDN Islands



# Pipeline-Based Load Balancing Design



# Performance Measurement: Experiments and Metrics

We will assess the design and implementation performance with the following experiments:

- ▶ Bidirectional “application-layer” and “transport-layer” communication across the gateway
- ▶ Unidirectional messages sent from IP and NDN hosts

We will collect the following metrics and model them as a function of the number of gateways  $n$  and estimated clients  $m$ :

- ▶ Unidirectional message translation overhead
- ▶ Unidirectional message trip time
- ▶ Bridge mode message latency (RTT)
- ▶ Bridge mode symmetric key establishment overhead time